

Paper B Frontier Visual Upgrade Report

Date: 20260619

Generated from frozen Paper B tables and newly downloaded frontier reference PDFs.

Figure 1. Source-anchored validation workflow




Table 1. Denominator-family extraction validity

Table 1. Denominator-family extraction validity

Table 1 | Paper B extraction validity by denominator family

Source: frozen 2,043-row M1-R scoring table. Scored-only accuracy is conditional on a scorable model answer; all-scorable accuracy treats abstentions as unresolved extraction work.


Denominator / claim role	Events / total	Correct / scored	Scored-only accuracy (95% CI)	All-scorable accuracy (95% CI)	Unresolved burden	Use in manuscript
Primary r validation pool	1,189 / 1,628	1,189 / 1,287	92.4% (90.8% to 93.7%)	73.0% (70.8% to 75.1%)	341 (20.9%)	Primary validation summary
Latent/source-flagged r	672 / 931	672 / 715	94.0% (92.0% to 95.5%)	72.2% (69.2% to 75.0%)	216 (23.2%)	Primary validation stratum
Direct/source-reported r	517 / 697	517 / 572	90.4% (87.7% to 92.5%)	74.2% (70.8% to 77.3%)	125 (17.9%)	Primary validation stratum
Converted beta/path	153 / 415	153 / 338	45.3% (40.0% to 50.6%)	36.9% (32.4% to 41.6%)	77 (18.6%)	Sensitivity and exception-handling stratum

Interpretation: the primary r strata support bounded extraction-validity claims; converted beta/path rows should remain a sensitivity/exception-handling stratum rather than a pooled accuracy numerator.

Figure 2. Risk-difference forest-style plot

Figure 2 | Accuracy difference against the converted beta/path sensitivity stratum

Risk differences use frozen Paper B denominator-family counts. Positive values mean the primary r stratum is more accurate than the converted beta/path stratum.



Note: This is not a vendor ranking. It tests whether directly source-supported r extraction behaves differently from converted/model-derived effect-size extraction.

Figure 3. Accuracy and review-burden profile

Figure 3 | Accuracy and review-burden profile

The x-axis is unresolved extraction burden, not elapsed time. A Lai-style time-savings plot requires reviewer-time logs that are not present in the current Paper B package.




Figure 4. MASEM claim gate

Figure 4 | MASEM identification gate

Covered construct pairs are not enough for a full target-model claim when complete-case studies collapse to zero.

